



University of
Pittsburgh

Algorithms and Data Structures 2

CS 1501



Fall 2022

Sherif Khattab

ksm73@pitt.edu

(Slides are adapted from Dr. Ramirez's and Dr. Farnan's CS1501 slides.)

Announcements

- Upcoming Deadlines
 - Homework 5: this Friday @ 11:59 pm
 - Lab 3: tonight @ 11:59 pm
 - Lab 4: next Monday @ 11:59 pm
 - Assignment 1: Monday Oct 10th @ 11:59 pm
- **Live support session** for Assignment 1
 - Recording and slides on Canvas
- **Student Support Hours** of the teaching team are posted on the Syllabus page

Previous lecture

- R-way Radix Search Tries
- De La Briandais (DLB) Tries

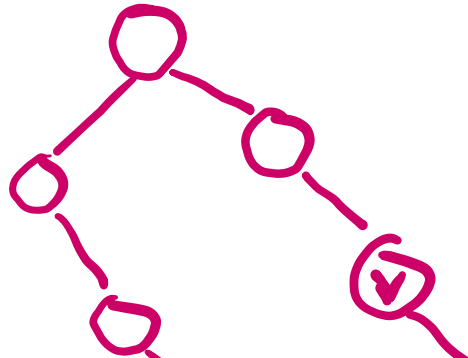
This Lecture

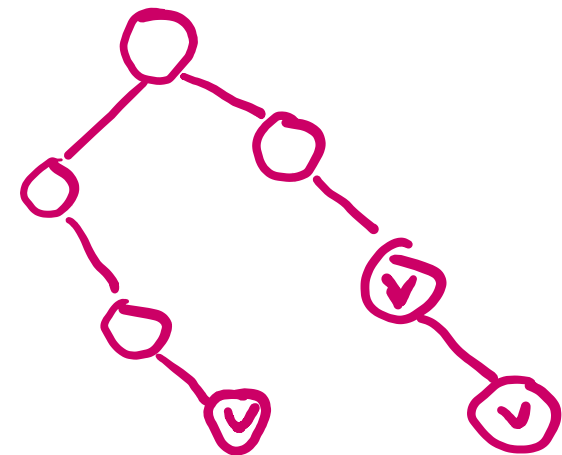
Compression

Muddiest Points

- **Q: When creating a DST, does it have to start handling keys starting with the leftmost bit, or can it also handle them by starting with the rightmost bit?**
- The algorithm can go either way on the bitstring of the key as long as the direction is the same for all operations

Muddiest Points

- **Q: Would a trie be able to contain a value with less bits than the root, and if so how?**
 - In a trie, none of the nodes (including the root) contains any key
 - If the question is “can a trie contain keys of different bit lengths?”,
 - the answer is yes
 - Interior nodes have non-null values in that case
 - The trie here has three keys
 - 011
 - 111
 - 11
- 



Muddiest Points


- **Q: How is a Trie different from a Red-Black BST?**
- a trie is different from a search tree because trie doesn't store the keys inside its nodes but a search tree does
- More in the next question

Muddiest Points

- **Q: when would you use a DST or an RST?**
- **Q: What's the application of RST?**
- DST and RST are efficient in checking if a target key is a prefix of any of the keys in the tree
 - e.g., making routing decisions in the Internet
- DSTs are preferred over BSTs when bits of keys are randomly distributed (i.e., the probability of each bit being zero is 0.5)
 - The DST will be balanced in this case without having to use the more complicated Red-Black BST
- RSTs are preferred over BSTs when bit lengths of keys are close to $\log n$
 - The RST will be balanced in this case without having to use the more complicated Red-Black BST
- Note that DST and RST don't provide the extra operations (e.g., predecessor and successor) provided by BST

Muddiest Points

- **Q: how can any node in 3's subtree replace 3 in DST example**
- Because all nodes in 3's subtree share a common prefix with length 1 with 3
 - The node that replaces 3 will still be found using the DST search algorithm
- For simplicity, we replace 3 with any **leaf** in its subtree



Muddiest Points

- **Q: Could you spend more time going through new lecture**
- Sometimes, addressing the muddiest points takes up a large portion of class time
- Usually, new material is embedded between the muddiest points responses

Muddiest Points

- **Q: Is the insertion position for DST based on the first bit that is different from the last insert? Or is it based on the relative comparison to last insert?**
- DST Add Algorithm for adding a key k and a corresponding value
 - if root is null, add k at the root and return
 - $\text{current} \leftarrow \text{root}$
 - if k is equal to the current node's key, replace value and return
 - if current bit of k is 0,
 - if left child is null, add k as left child
 - else continue to left child
 - if current bit of k is 1,
 - if left child is null, add k as right child
 - else continue to right child

Muddiest Points

- **Q: When is DST preferable to radix search trie?**
- **A: When bit lengths of keys are $\gg \log n$**

Muddiest Points

- **Q: I don't understand the advantage of making another node in the DLB instead of the tree structure.**
- **A: DLB saves space when the number of children per node in an R-way RST is small**

Muddiest Points

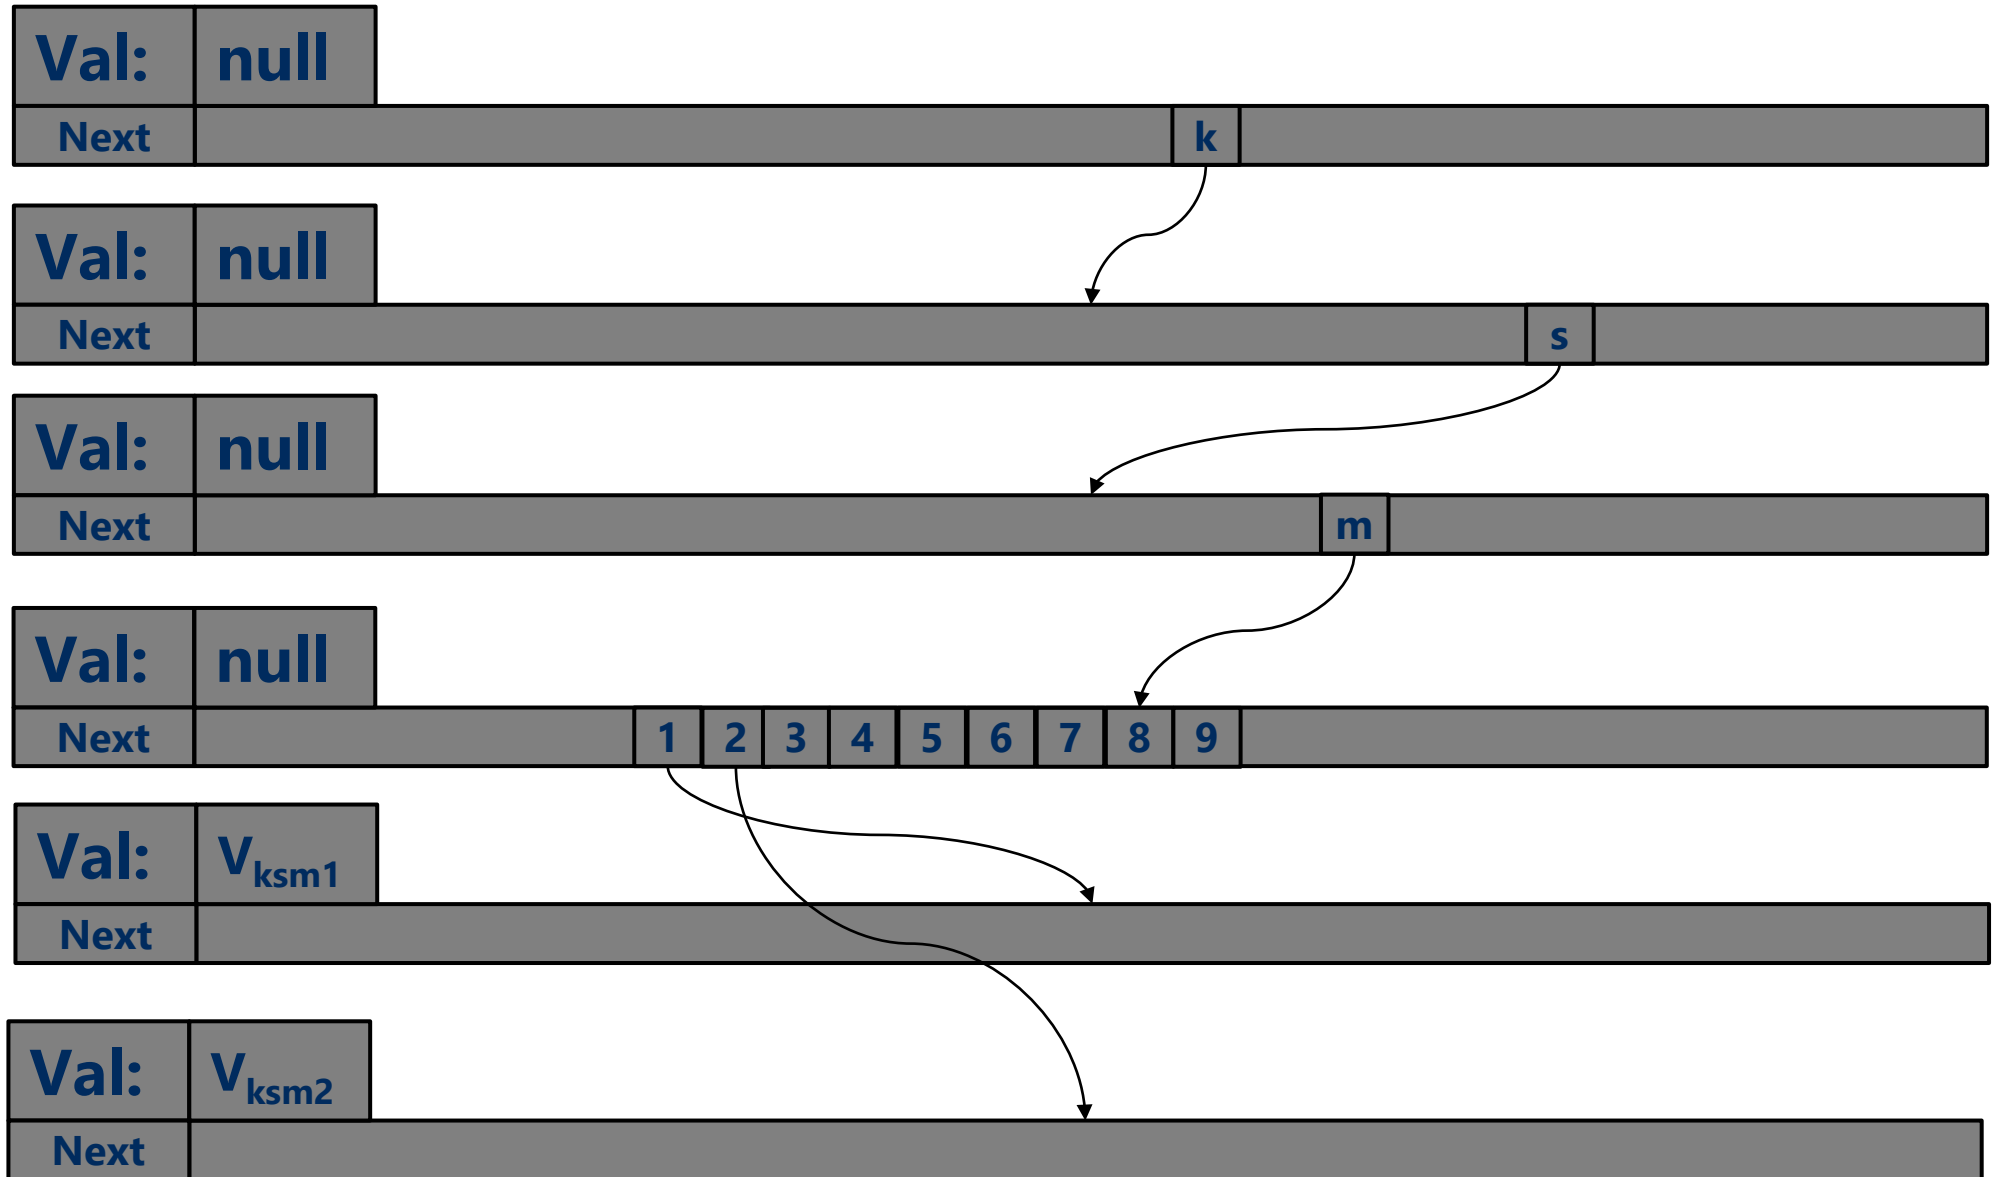
- **Q: How does DLB save space over r way trie?
Example please?**
- Let's the set of keys:
 - ksm1 ... ksm9
- How big does an 256-way RST take vs. a DLB trie?

R-way RST

```
private class Node {  
    private Object val;  
    private Node[] next;  
  
    private Node(){  
        next = new Node[R];  
    }  
}
```

Each node takes $4 \cdot (R+1) = 4 \cdot 257 = 1028$ bytes,
assuming 4 bytes per reference variable

R-way RST



R-way RST

We will end up with $4 + 9 = 13$ nodes

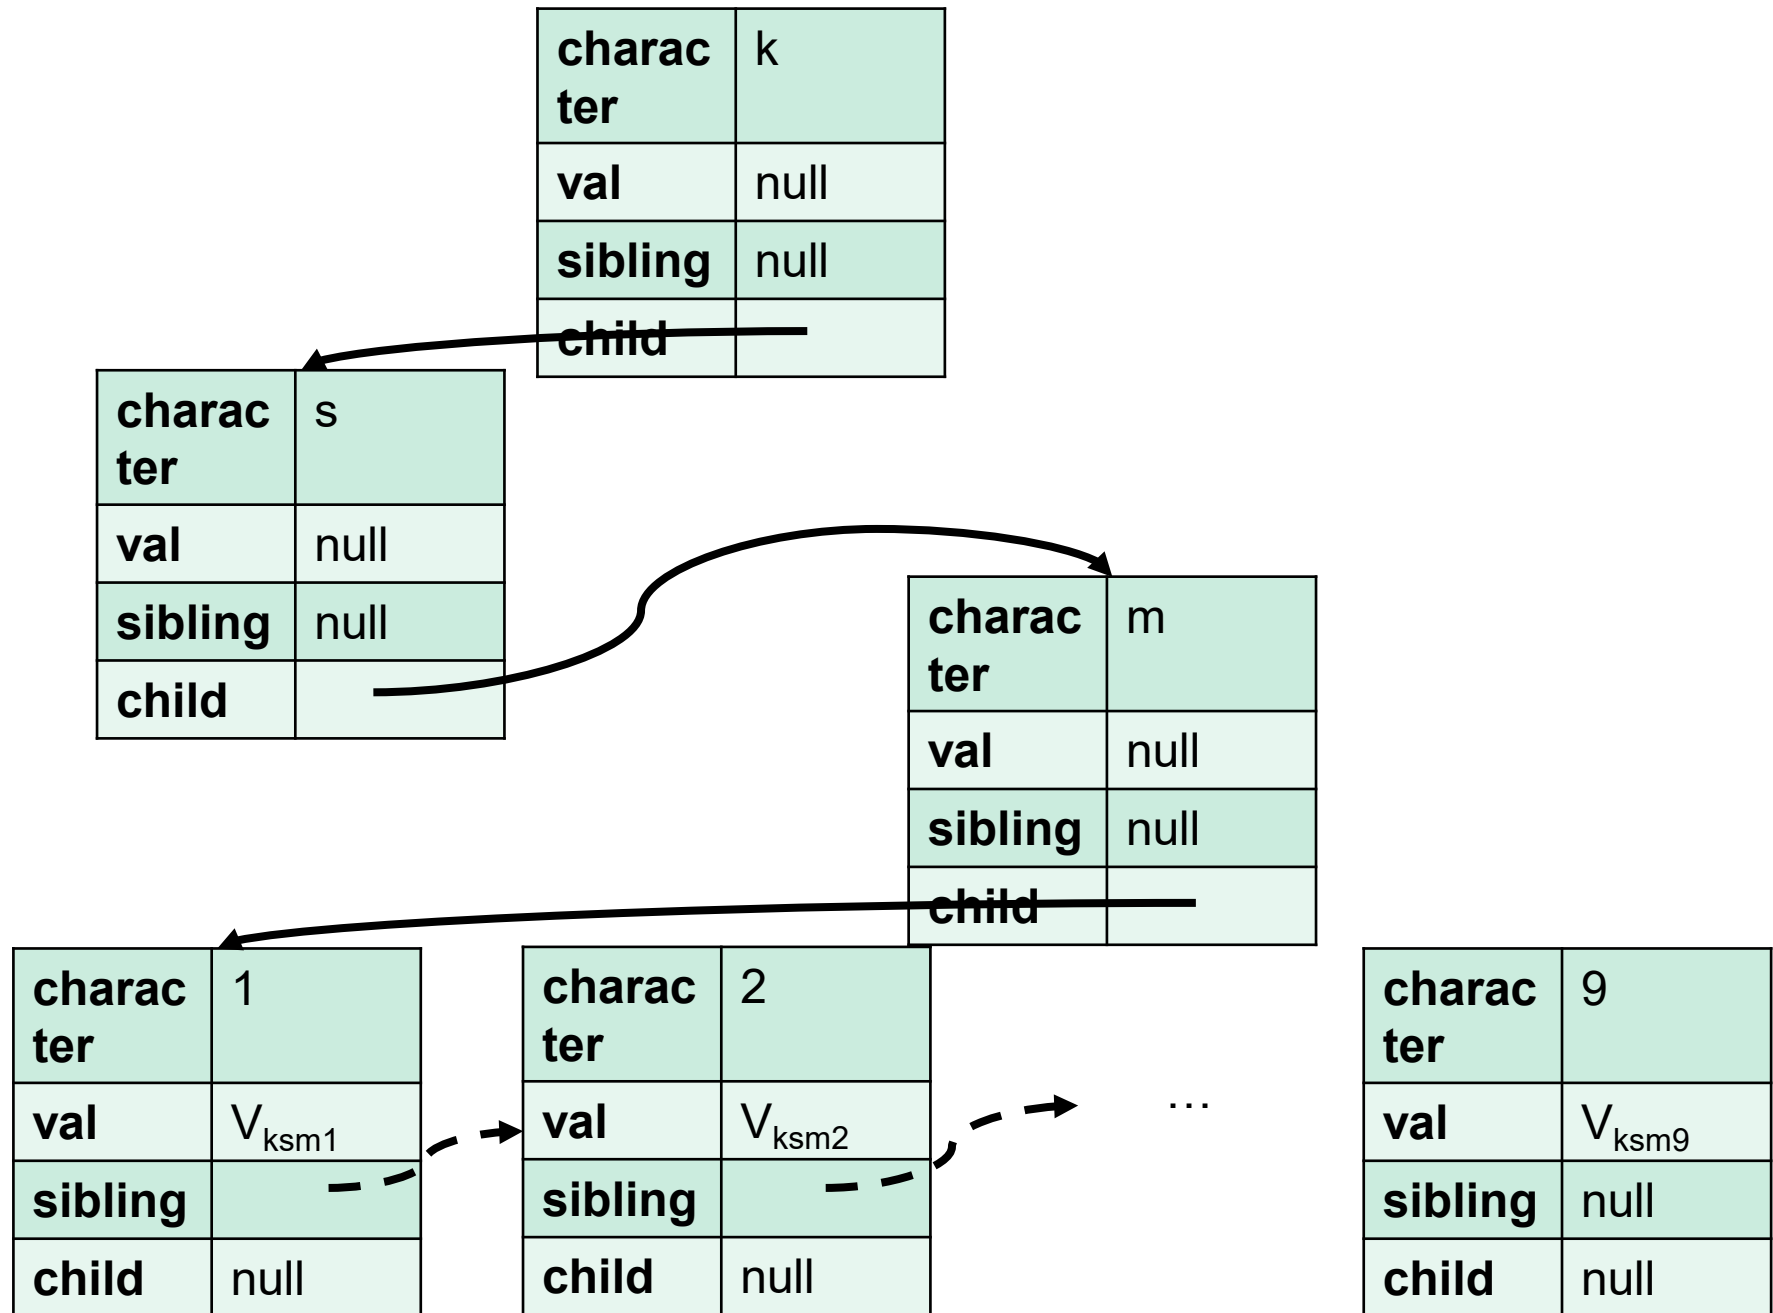
Total space is $13 * 1028 = 13,364$ bytes

DLB Trie

```
private class DLBNode<T> {  
    private Character character;  
    private Object val;  
    private Node sibling;  
    private Node child;  
}
```

Each node takes $4 \times 4 = 16$ bytes, assuming 4 bytes per reference variable

DLB Trie



DLB Trie

We will end up with 12 nodes

Total space is $12 * 16 = 192$ bytes

Compare to 13,364 bytes with an R-way RST

Muddiest Points

- **Q: What determines the number of bits you use for the bit representation of a key in DSTs and RSTs?**
- Typically, the number of bits is determined by the application
 - e.g., keys are Pitt usernames, PeopleSoft IDs, English sentences, etc.
- It is better to re-encode the keys to have a bit length of $\log n$ bits each
 - Requires extra space to store the mapping from old keys to new keys
 - sometimes not possible: e.g., when n is not known in advance
- Better yet, we can assign bit lengths based on frequency of access:
 - Shorter bitstrings for more frequently accessed keys
 - Results in smaller **average** search time

$$\text{average Case runtime} = \sum_{\text{all cases}} P(\text{Case}_i) \times \text{runtime for Case}_i$$

Muddiest Points

- **Q: Not really a muddiest point, but it would be extremely helpful to see actual code (not pseudo code) next to some of these trees**
- You will see code in the recitations

Muddiest Points

- **Q: DLB do what?**
- De La Briandais (DLB) Trie
 - tree-like structure used for searching when keys are sequences of characters
 - each nodelet
 - stores one character,
 - points to a sibling (linked list of siblings), and
 - points to a child
 - worst-case running time is $O(wR)$
 - w : number of characters in the key
 - R : alphabet size
 - worst-case can be avoided by using DLB only when the sibling lists are short
 - check add algorithm in previous lecture

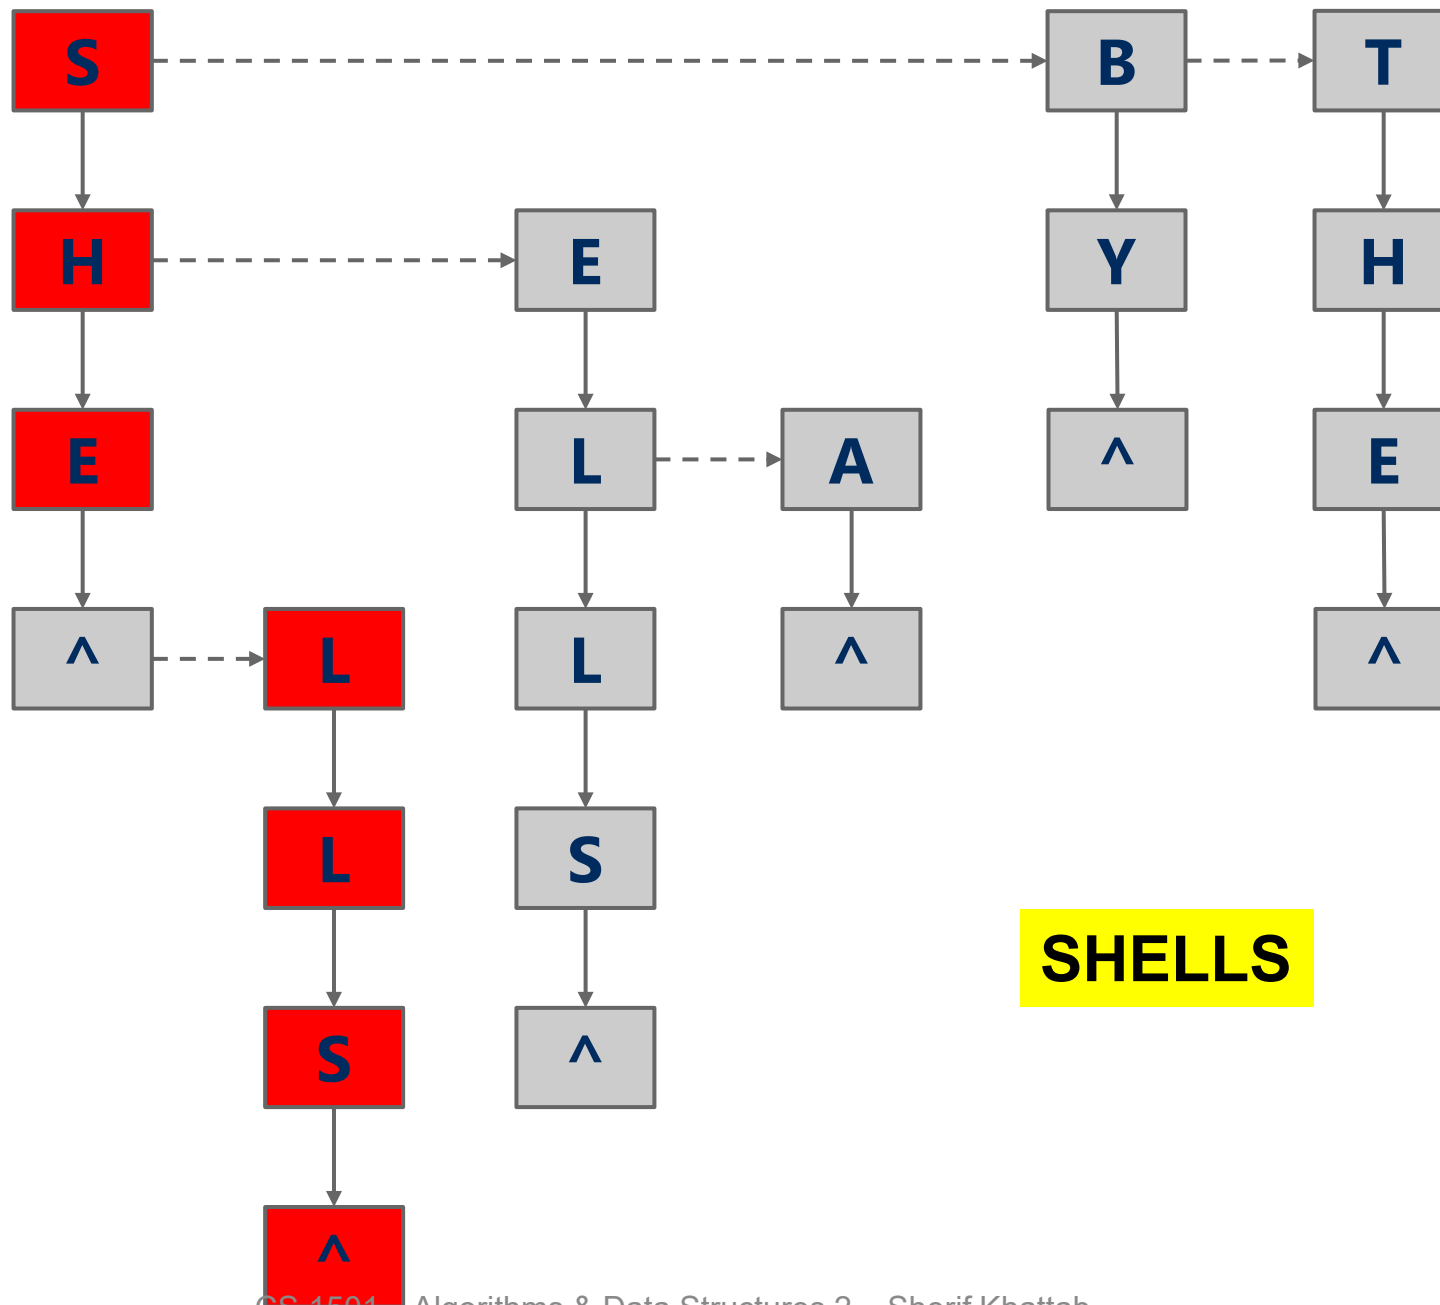
Muddiest Points

- **Q: Just wanted to quickly double check some details about DSTs: The max height of a DST is the number of key bits + 1**
- Correct
- **Q: and the max comparisons you can do is the key's bit length, correct?**
- No, it is also $b + 1$

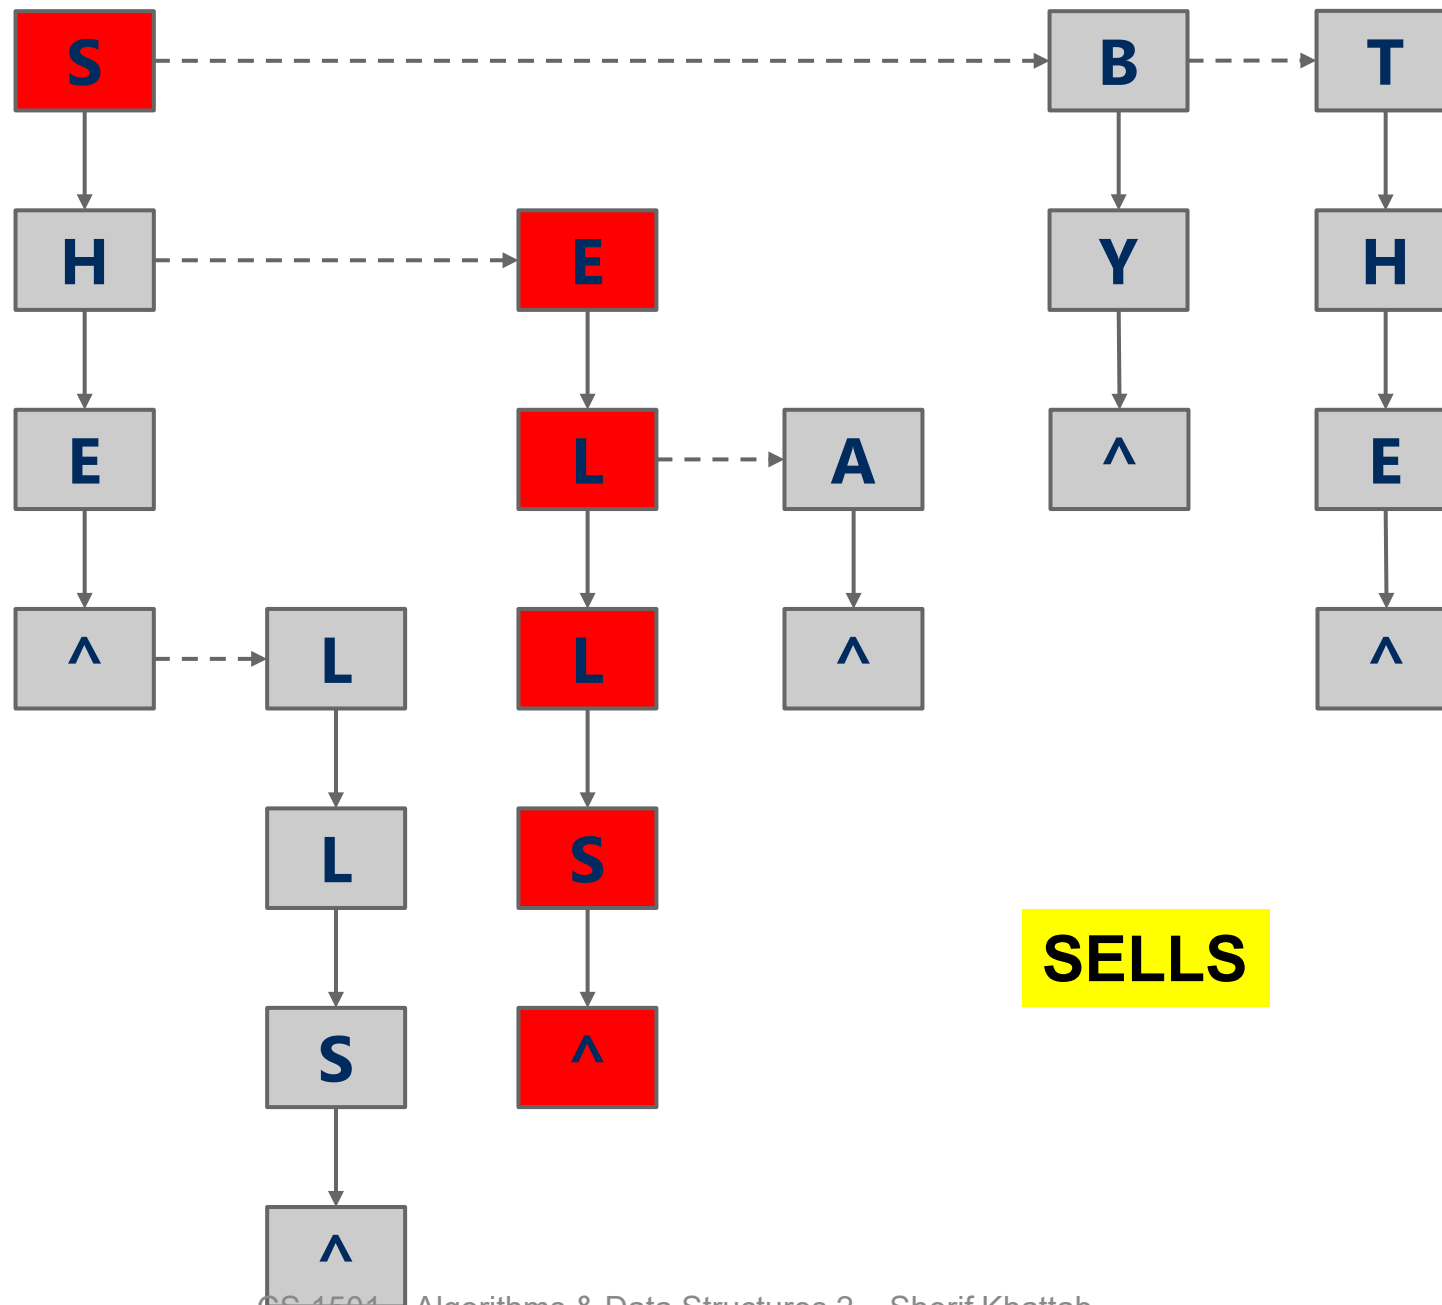
Muddiest Points

- **Q: why was there two paths for "shell" on the DLB example?**
- There was only one path!

DLB Example



DLB Example



Muddiest Points

- **Q: What is the point of the sentinel character? How does that implementation of the DLB differentiate it from the other DLB implementation?**
- If the DLB stores only keys (without corresponding values), we don't need the val field in the DLBNode
- But, val helped us determine if the node we stop at corresponds to a key or not
 - when val is not null, the node corresponds to a key
- Without val, we need a different method: using a sentinel
 - the sentinel is added to each key before adding and before searching
 - a key is found when the key with the sentinel is found
 - e.g., adding she results in adding she^
 - searching for she becomes searching for she^
 - if "she^" is found then she is a key
 - if only "she" is found (without the sentinel), she is not a key

Muddiest Points

- **Q: what sorting methods have the data explicitly and when is it implicit.**
- Tries store keys implicitly, whereas trees store keys explicitly inside tree nodes
- Check node structure for a tree vs. a trie

Muddiest Points

- **Q: Can you re-explain what you meant by $w = b/\text{ceiling}(\log R)$?**
- The string “she” has 3 characters ($w=3$)
- If we look at the bit representation of “she”, assuming each character is an extended ASCII character (i.e., 8-bit character), the number of bits will be $b = 3 * 8 = 24$
- For extended ASCII, the alphabet size is $R = 2^8 = 256$
- $w = b/8 = b/\log R$

Problem of the Day: Compression

- Input: A file containing a sequence of characters
 - n characters
 - each encoded as an 8-bit Extended ASCII
 - total file size = $8*n$
- Output: A shorter bitstring
 - of length $< 8*n$
 - such that the original sequence can be fully restored from the bitstring

ASCII Encoding

A set of R symbols can be represented using fixed-size encoding of length k bits each

- iff $2^k \geq R$
- that is, $k = \lceil \log R \rceil$

